

# CS 503 - Fall 2017

## Lab 1: Process Scheduling (100 pts)

**Due Date: Thursday, Sept. 28, 2017, 11:59PM**

---

In this lab you will implement a two-level process scheduler and few scheduling policies in XINU that will avoid *starvation* of processes and improve load balancing between processes. At the end of this lab you will be able to explain the advantages and disadvantages of the scheduling policies implemented and evaluated in XINU.

Starvation is produced in XINU when we have two or more processes eligible for execution that have different priorities. The scheduling invariant in XINU assumes that, at any time, the highest priority process eligible for CPU service is executing, with round-robin scheduling for processes of equal priority. Under this scheduling policy, processes with the highest priority, if ready, will always be executing. As a result, processes with lower priority may never receive CPU time. For ease of discussion we call "the set of processes in the ready list and the current process" as *eligible* processes.

All processes in XINU will be put into one of the two scheduling groups: proportional share group and TS group. Within each group, processes will be scheduled via different scheduling policy. When *resched()* is invoked, it should decide which group should occupy the CPU at first. Then it picks up a process from this group to run. In this lab, scheduler should run **Aging Scheduler** to pick the group. After the group is decided, it should pick up one process in this group via group-specific scheduling policy. Proportional share group should apply **Proportional Share Scheduler**; TS group should apply **TS Scheduler**. In the following we will introduce the three scheduling policies.

---

### 1. Aging Scheduler [20pts]

The scheduling policy for process group scheduling is **Aging scheduler**. On each rescheduling operation, the scheduler should count the number of processes from different groups in ready queue (e.g. there are 3 processes from proportional group and 4 processes from TS group in ready queue). It should increase the priority of the groups by their number of processes (priority of proportional group increases by 3 and priority of TS group increases by 4). Then it should pick up the group with highest priority. If there are processes only from one group, it just picks up this group directly.

#### Implementation sketch:

Each group has an *initial priority*. By default the priority of those two groups are both 10. And it could be changed via the call to *chgprio()*. Every time the scheduler is called it takes the following steps.

- If the *current* process belongs to group A, the priority of group A is set to the *initial priority* assigned to it.
- The priorities of all groups are incremented by the number of processes (in that group) in the ready queue. Here you don't count the *current* process and *null* process

Note that during each call to the scheduler, the complete ready list has to be traversed. Also, when both groups have the same priority, choose Proportional Share scheduling policy.

---

## 2. Proportional Share Scheduler [40pts]

Every process in the proportional scheduling group has a scheduling parameter called *rate*. In the following, we will assume that all the processes in question belong to the proportional scheduling group. For a process  $i$  let us denote the *rate* as  $R_i$ . Every process  $i$  has a *priority value*  $P_i$ . Initially, all the processes start with a *priority value* 0 ( $P_i = 0$ ). Whenever a rescheduling occurs, the *priority value*  $P_i$  of the **currently** running process is updated as follows

$$P_i := P_i + (t * 100 / R_i)$$

where  $t$  is the CPU ticks consumed by the process since it was last scheduled.  $R_i$  is a percentage value and takes values between 1 and 100.

Now the scheduler schedules an eligible process with the **smallest**  $P_i$ . Whenever a process is scheduled the first time or is scheduled after blocking, its  $P_i$  value is updated as

$$P_i := \max ( P_i, T )$$

where  $T$  is the number of CPU ticks that have elapsed since the start of the system.

Intuitively, you can think of this policy as one that gives the processes some guarantees about their CPU share. If a process has a rate  $R_i$ , then it is guaranteed at least  $R_i$  percent of CPU time, provided the sum of all  $R_i$  values is less than 100. As the CPU usage of a process increases, its  $P_i$  value also increases depending on its  $R_i$ . If you have a large  $R_i$ , then your  $P_i$  increases more slowly and hence giving you a larger share of the CPU. Thus,  $R_i$  can be considered as a *share* of the CPU for the process  $i$ .

The second formula can be intuitively understood from the following example: Consider two processes A, B starting at time 0 and running continuously with rate 50 and 40 respectively. Let us say that another process C is created after 100,000 ticks with rate 10. C will start executing with a *priority value* of 0 (if the second formula were not to be applied) and hence will hog the CPU for a very long time and processes A, B have to wait for long to get the CPU back and would not enjoy their share of the CPU till all the  $P_i$  values level off. On the other hand, if the process C starts with a *priority value* 100,000 instead of 0, then it will run only for a short amount of time before yielding the CPU back to A and B.

### Implementation sketch:

According to this policy, processes are scheduled in increasing order of their priority, i.e., a process with the lowest *priority value*  $P_i$  will be scheduled first. However, note XINU works in the opposite way, i.e., a process with the highest priority is scheduled first. In order to overcome this mismatch, we

can maintain an internal variable  $P_i$  for every process which will contain the *priority value*. Let us denote the XINU process priority of a process  $i$  to be  $PRIO_i$ . Then  $PRIO_i$  can be calculated as  $PRIO_i = MAXINT - P_i$ . As  $P_i$  increases,  $PRIO_i$  decreases and the process will get lesser share of CPU.

To summarize, at every reschedule operation, the proportional share scheduler does the following:

- The  $P_i$  value of the current process is modified and its  $PRIO_i$  value updated.
- The scheduler chooses for execution the process with the highest priority, choosing from the processes in the ready list *and* the current process.

Also, whenever a process is scheduled for the first time or immediately after blocking, then the  $P_i$  value is changed as indicated above and the  $PRIO_i$  is updated to reflect the change. You need to identify when and how to change  $P_i$ . Other than using previous method, you can keep track of minimum of  $P_i$ , which requires extra traversals in readylist.

### Benchmark sketch:

To evaluate Proportional Share Scheduling, you will need to experiment with two situations : (i) when  $R_i$  has the equal rate. (ii) when  $R_i$  has the different rates. Also, you should consider when some processes in this group are executed late or are blocked and come back later.

---

## 3. TS Scheduler [40pts]

A TS (timeshare) scheduler attempts to classify processes into CPU-bound and I/O-bound categories such that I/O-bound processes can be assigned higher priority for increased responsiveness without sacrificing fairness (i.e., starving CPU-bound processes). This is so since I/O-bound processes tend to voluntarily relinquish the CPU by issuing blocking system calls before their time quanta has expired.

### Implementation sketch:

Implementing a TS scheduler entails two aspects: first, identification as I/O- or CPU-bound, and second, adapting process priorities and time quantum. For the first step, we will use a simple (but also efficient) criterion, namely, whether the current process (if it belongs to the TSSCHED class), at the moment *resched()* is called, voluntarily relinquished the CPU or was forcibly interrupted by a clock interrupt handler because its time quantum expired. This may be determined inside *resched()* by checking the global variable *preempt*. If a TSSCHED class process voluntarily relinquished CPU, it is considered I/O-bounded. Otherwise, a process is classified as CPU-bounded process (note that this only takes into account the most recent history). After classification, priority and time quantum will be dynamically adjusted based on the classification. The dynamic adjustment will be based on Solaris [UNIX dispatch table](#). When a process is viewed as IO-bounded process, it gets a higher priority (*ts\_slpret*) but less time quantum (corresponding *ts\_quantum*). If it is considered as CPU-bounded process, it gets a lower priority (*ts\_tqexp*) and more time quantum (corresponding *ts\_quantum*). In UNIX dispatch table, please ignore *ts\_maxwait* and *ts\_lwait* columns.

## Benchmark sketch:

To evaluate how well the TS scheduling implementation in XINU balances fairness and performance of CPU- and I/O-bound processes, consider the following three test case scenarios.

1. *All processes are CPU-bound.* In the first test case, create 6 processes that run the same program *cpubound()*. Put the code of *cpubound()* in a separate file *cpubound.c*. Assign the same initial priority value 1 when calling *create()* from *main()*. The code structure of *cpubound()* should follow:

```
for (i=0; i<LOOP1; i++) {
    for (j=0; j<LOOP2; j++) {
        // Insert code that performs memory copy (slow) and/or
        // ALU operations (fast).
        // Note: this loop consumes significant CPU cycles.
    }
    // Using kprintf print the pid followed the outer loop count i,
    // the process's priority and remaining time slice (preempt).
}
```

You will need to experiment with different values of *LOOP1* and *LOOP2* to induce CPU sharing via context switching and resultant output that can be easily interpreted for gauging your implementation of TS scheduling performance. Your runs should be long enough so that you can conclude from the CPU share output that the system has reached a "steady state."

2. *All processes are I/O-bound.* Follow the same steps as above but replace *cpubound()* by *iobound()* (in *iobound.c*) which has the same code structure as *cpubound()* except that the code in the inner loop (for ALU and memory copy) is replaced by a single call to *sleepms()*. By varying the argument of *sleepms()* you can modulate the degree to which *iobound()* is prone to block and voluntarily relinquish the CPU. When testing, use the same sleep time as argument to *sleepms()* for all 6 instances of I/O-bound processes.
3. *Half-and-half:* . Create 6 processes where half execute *cpubound()* and the other half execute *iobound()*. In the first part of the evaluation, under this mixed workload of CPU- and I/O-bound processes, determine if the 3 CPU-bound processes --- among themselves --- achieve equal sharing of CPU cycles as indicated by the output. Do the same for the 3 I/O-bound processes with their sleep time arguments to *sleepms()* fixed to the same value. Evaluate CPU sharing between the two groups of processes --- CPU- and I/O-bound

**NOTE: You don't need to submit *cpubound.c* and *iobound.c***

---

## 4. System call implementation

**create(void \*funcaddr, uint32 ssize, int group, pri16 priority, ... )**

Please add a new argument, *group* to this function (before *priority*) to this function. *group* should be either PROPORTIONALSHARE or TSSCHED. And for processes of proportional group, *priority* is used to define  $R_i$ . For processes which created by XINU by default (e.g. main, null), you can put them into either PROPORTIONALSHARE or TSSCHED.

## **chgprio(int group, pri16 newprio)**

You should add this new system call to change the initial priority of groups. Here you could have a look at system call *chprio* which changes processes' priority. The new group priority will be effective from the next scheduling.

## **resched()**

This system call will be invoked for scheduling a process to run. Most of your work will be done here. So please understand each line of code of this system call before you start to implement. **One thing to note** is that this lab's focus is not in scheduling's efficiency, but its correctness. Therefore, you can share readylist for both scheduling policy.

---

## **5. Turnin Instructions**

Turnin instructions for Lab1 code (electronic):

1. Make sure to turn off debugging output. Also, please do not rename xinu-fall2017, or any of its subdirectories.
2. Go to *xinu-fall2017/compile* directory and do "make clean".
3. Go to the directory of which your *xinu-fall2017* directory is a subdirectory. (eg) if */homes/XXX/xinu-fall2017* is your directory structure, goto */homes/XXX*
4. Submit using the following command:  
***turnin -c cs503 -p lab1 xinu-fall2017***
5. After submitting your files, you may want to check that it's indeed submitted using command:  
***turnin -v -c cs503 -p lab1***

***You can write code in main.c to test your procedures, but please note that when we test your programs we will replace the main.c file. Therefore, do not put any functionality in the main.c file.***

---